

Profiling and Complexity

David Croft

Coventry University

david.croft@coventry.ac.uk

2017

Overview

1 Profiling

- Efficiency
- Optimization
- Profilers

2 $O()$ notation

- Simple algorithms
- Good algorithms
- Bad algorithms

3 Recap

When writing software think about its efficiency.

- Time.
- Memory.

When writing software think about its efficiency.

- Time.
- Memory.
- Time vs Memory.
 - Can you trade one for the other
 - I.e. data stored in RAM costs memory but saves time.
 - I.e. data stored on hard drive saves memory but costs time.

When writing software think about its efficiency.

- Time.
- Memory.
- Time vs Memory.
 - Can you trade one for the other
 - I.e. data stored in RAM costs memory but saves time.
 - I.e. data stored on hard drive saves memory but costs time.
- Optimization makes software run faster/leaner/better.

"Premature optimization is the root of all evil"

–Donald Knuth

"Premature optimization is the root of all evil"

–Donald Knuth

For any large piece of code you should:

"Premature optimization is the root of all evil"

–Donald Knuth

For any large piece of code you should:

- Write clear, easily understood code. Focus on getting the behaviour right, not on performance.

"Premature optimization is the root of all evil"

–Donald Knuth

For any large piece of code you should:

- Write clear, easily understood code. Focus on getting the behaviour right, not on performance.
- Test the performance.
 - It may be fine.

"Premature optimization is the root of all evil"

–Donald Knuth

For any large piece of code you should:

- Write clear, easily understood code. Focus on getting the behaviour right, not on performance.
- Test the performance.
 - It may be fine.
- Profile your code to get the baseline performance.
 - So that you know if you are making things better or worse.

"Premature optimization is the root of all evil"

–Donald Knuth

For any large piece of code you should:

- Write clear, easily understood code. Focus on getting the behaviour right, not on performance.
- Test the performance.
 - It may be fine.
- Profile your code to get the baseline performance.
 - So that you know if you are making things better or worse.
- Focus your efforts on the code that is consuming all the time.
 - E.g. small pieces of code that get called multiple times.

Profiling is a method of analysing your code to identify the impact of the different functions/classes/sections etc.

Instrumentation profilers

- Add extra bits of code to track time/memory/function calls.
 - Can be done manually.
 - But automatic is better.
- Accurate.
 - But slows things down.

Statistical profilers

- Regularly checks the software state.
- Accurate-ish.
 - Based on statistical sampling.
 - Doesn't slow things down.

In this example which function takes the most time?

- `fast_math_function()` or `slow_math_function()`?

Example

1

```
def fast_math_function(a, b):  
    time.sleep(0.00001)  
    return a + b  
  
def slow_math_function(a, b):  
    time.sleep(3)  
    return a + b  
  
def main():  
    for i in range(int(1.0000)):  
        slow_math_function(42, 69)  
  
    for i in range(int(100000)):  
        fast_math_function(42, 69)  
  
if __name__ == '__main__':  
    sys.exit(main())
```

lec_functions.py

In this example which function takes the most time?

- `fast_math_function()` or `slow_math_function()`?
- Why don't we just profile it and find out?

Example

1

```
def fast_math_function(a, b):  
    time.sleep(0.00001)  
    return a + b  
  
def slow_math_function(a, b):  
    time.sleep(3)  
    return a + b  
  
def main():  
    for i in range(int(1.0000)):  
        slow_math_function(42, 69)  
  
    for i in range(int(100000)):  
        fast_math_function(42,69)  
  
if __name__ == '__main__':  
    sys.exit(main())
```

lec_functions.py

```
» python3 -m cProfile lec_functions.py
```

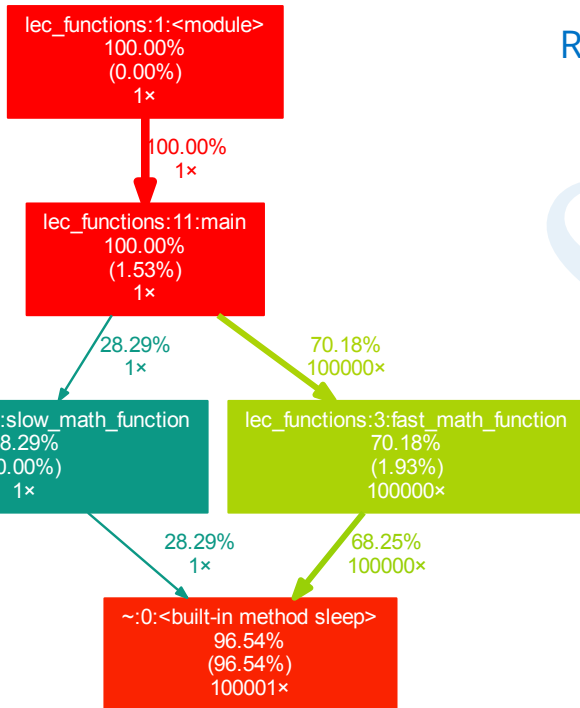
```
200007 function calls in 10.362 seconds
```

```
Ordered by: standard name
```

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	10.362	10.362	lec_functions.py:1(<module>)
1	0.137	0.137	10.362	10.362	lec_functions.py:11(main)
100000	0.171	0.000	7.222	0.000	lec_functions.py:3(fast_math_function)
1	0.000	0.000	3.003	3.003	lec_functions.py:7(slow_math_function)
1	0.000	0.000	10.362	10.362	{built-in method exec}
1	0.000	0.000	0.000	0.000	{built-in method exit}
100001	10.054	0.000	10.054	0.000	{built-in method sleep}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' object}

Things to note:

- Total time - time spent in each function.
- Cumulative time - time spent in each function AND the functions it calls.



Results visualised

C

Results passed through
Graphviz/gprof2dot.

- A profiling visualisation tool.

Profiling is very useful in determining the actual performance of your code.

- Unexpected bottlenecks.
- Problems in 3rd party libraries etc.
- Not so good at measuring how code will scale.
 - Change in response to different inputs.
- Algorithmic complexity.
- Certain algorithms are known to be better than other algorithms.

Used to describe complexity in terms of time and/or space.


- Commonly encountered examples...
 - $O(1)$, $O(\log n)$, $O(n)$, $O(n \log n)$, $O(n^2)$,
 $O(2^n)$ and $O(n!)$
- n refers to the size of the problem.
 - E.g. n values to be sorted.
 - E.g. n values to be searched.
- $O()$ notation describes the worst case scenario.
 - Usually, unless otherwise stated.
- Main idea is to capture the dominant term: the thing that is most important when the size of the input (n) gets big.

Linear complexity.

- n is directly proportional to time/space required
 - E.g. n doubles then time/space doubles.
- E.g. linear/sequential search.

```
a = [ 0, 1, 2, 3, 4, 5, 6, 7, 42 ]
```

```
for i in a:  
    if i == 42:  
        print('Found it')  
        break
```




Linear complexity.

- n is directly proportional to time/space required
 - E.g. n doubles then time/space doubles.
- E.g. linear/sequential search.

```
a = [ 0, 1, 2, 3, 4, 5, 6, 7, 42 ]
```

```
for i in a:  
    if i == 42:  
        print('Found it')  
        break
```




Linear complexity.

- n is directly proportional to time/space required
 - E.g. n doubles then time/space doubles.
- E.g. linear/sequential search.

```
a = [ 0, 1, 2, 3, 4, 5, 6, 7, 42 ]
```

```
for i in a:  
    if i == 42:  
        print('Found it')  
        break
```

(n)
(n)




Linear complexity.

- n is directly proportional to time/space required
 - E.g. n doubles then time/space doubles.
- E.g. linear/sequential search.

```
a = [ 0, 1, 2, 3, 4, 5, 6, 7, 42 ]
```

```
for i in a:  
    if i == 42:  
        print('Found it')  
        break
```

(n)
(n)
(1)




Linear complexity.

- n is directly proportional to time/space required
 - E.g. n doubles then time/space doubles.
- E.g. linear/sequential search.

```
a = [ 0, 1, 2, 3, 4, 5, 6, 7, 42 ]
```

```
for i in a:  
    if i == 42:  
        print('Found it')  
        break
```

(n)
(n)
(1)
(1)




Linear complexity.

- n is directly proportional to time/space required
 - E.g. n doubles then time/space doubles.
- E.g. linear/sequential search.

```
a = [ 0, 1, 2, 3, 4, 5, 6, 7, 42 ]
```

```
for i in a:
    if i == 42:
        print('Found it')
        break
```

(n)
(n)
(1)
(1)




- So the algorithm takes $n + n + 1 + 1 = 2n + 2$ operations.
 - BUT! We would say it has complexity $O(n)$, constant values are irrelevant.

Constant complexity.

- n doesn't matter.
- Always takes same time/space.
- E.g. getting first item in an array.

```
a = [ i for i in range(100) ]  
b = [ i for i in range(1000000) ]
```

```
print(a[0])  
print(b[0])
```




Constant complexity.

- n doesn't matter.
- Always takes same time/space.
- E.g. getting first item in an array.

```
a = [ i for i in range(100) ]  
b = [ i for i in range(1000000) ]
```

```
print(a[0])  
print(b[0])
```

(1)



Constant complexity.


- n doesn't matter.
- Always takes same time/space.
- E.g. getting first item in an array.

```
a = [ i for i in range(100) ]  
b = [ i for i in range(1000000) ]
```

```
print(a[0])  
print(b[0])
```

(1)

(1)




Constant complexity.

- n doesn't matter.
- Always takes same time/space.
- E.g. getting first item in an array.

```
a = [ i for i in range(100) ]      (n)
b = [ i for i in range(1000000) ] (m)

print(a[0])                       (1)
print(b[0])                       (1)
```




Quadratic complexity.

- A lot of simple sorting algorithms are $O(n^2)$.
- Nested **for** loops are common example.
- $O(n^3)$, $O(n^4)$, $O(n^m)$ etc. are all possible.
- Polynomial time.

```
print('The n times tables')
```

```
for i in range(n):  
    for j in range(n):  
        print(i*j)
```




Quadratic complexity.

- A lot of simple sorting algorithms are $O(n^2)$.
- Nested **for** loops are common example.
- $O(n^3)$, $O(n^4)$, $O(n^m)$ etc. are all possible.
- Polynomial time.

```
print('The n times tables') (1)
```

```
for i in range(n):  
    for j in range(n):  
        print(i*j)
```




Quadratic complexity.

- A lot of simple sorting algorithms are $O(n^2)$.
- Nested **for** loops are common example.
- $O(n^3)$, $O(n^4)$, $O(n^m)$ etc. are all possible.
- Polynomial time.

```
print('The n times tables') (1)
```

```
for i in range(n):           (n)  
    for j in range(n):  
        print(i*j)
```




Quadratic complexity.

- A lot of simple sorting algorithms are $O(n^2)$.
- Nested **for** loops are common example.
- $O(n^3)$, $O(n^4)$, $O(n^m)$ etc. are all possible.
- Polynomial time.

```
print('The n times tables') (1)
```

```
for i in range(n):           (n)  
    for j in range(n):       (n*n)  
        print(i*j)
```




Quadratic complexity.

- A lot of simple sorting algorithms are $O(n^2)$.
- Nested **for** loops are common example.
- $O(n^3)$, $O(n^4)$, $O(n^m)$ etc. are all possible.
- Polynomial time.

```
print('The n times tables')    (1)

for i in range(n):            (n)
    for j in range(n):        (n*n)
        print(i*j)            (n*n)
```


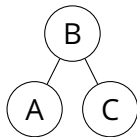


$O(\log n)$

1

Logarithmic complexity.

- Bit more complicated.
- Think binary search.





$O(\log n)$

1


Logarithmic complexity.

- Bit more complicated.
- Think binary search.




Logarithmic complexity.

- Bit more complicated.
- Think binary search.




$$n = 3$$

$$O(\log n) = 1.58 \Rightarrow 1$$




Logarithmic complexity.

- Bit more complicated.
- Think binary search.



$$n = 3$$


$$O(\log n) = 1.58 \Rightarrow 1$$



$$n = 7$$


Logarithmic complexity.

- Bit more complicated.
- Think binary search.



$$n = 3$$

$$O(\log n) = 1.58 \Rightarrow 1$$




$$n = 7$$

$$O() = 2.81 \Rightarrow 2$$


$O(\log n)$ complexity.

- Rate of increase gets lower and lower.
- $\log_2(100)$ is only 6.
- $\log_2(1000000000000)$ (trillion) is only 39.



$$n = 15$$

$$O() = 3.91 \Rightarrow 3$$

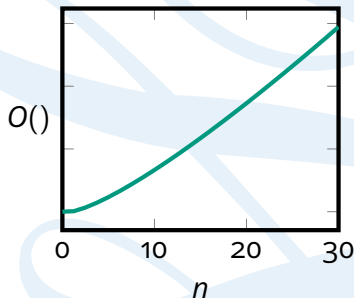


$O(n \log n)$

A


Loglinear complexity.

- Looks more difficult than it is.
- $O(n \log n)$ means, do $O(\log n)$ n times.
- E.g. binary search for n items.
 - Binary search is $O(\log n)$.
 - Doing n binary searches.
 - So $O(n \log n)$.
- Lots of good sorting algorithms are $O(n \log n)$.




Exponential complexity.

- Very, very bad.
- Each additional value doubles the time/space.
- Doesn't scale.
- $O(3^n)$, $O(4^n)$ etc. are all possible.




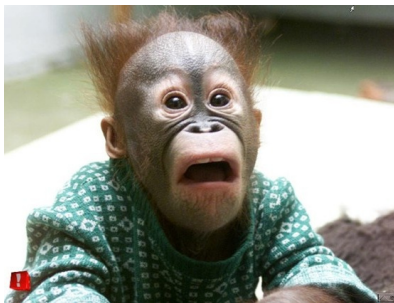
Factorial complexity.

- Just awful.
- Every possible combination of n items.
- Brute force travelling salesman is $O(n!)$.
- Totally impractical even for small values of n .




Factorial complexity.

- Just awful.
- Every possible combination of n items.
- Brute force travelling salesman is $O(n!)$.
- Totally impractical even for small values of n .



Different $O()$ == wildly different complexity.

		n		
		2	10	100
Best	$O(1)$	1	1	1
	$O(\log n)$	1	3	6
	\uparrow $O(n)$	2	10	100
	\downarrow $O(n \log n)$	2	33	664
	$O(n^2)$	4	100	10000
Worst	$O(2^n)$	4	1024	$1.27 \cdot 10^{30}$
	$O(n!)$	2	3628800	$9.33 \cdot 10^{157}$





Complexity isn't the same as efficiency.

- A good $O(n^2)$ implementation can be better than a bad $O(n)$.
 - For a while.
- Eventually, as n increases, $O(n)$ will always outperform $O(n^2)$ etc.

Complexity vs. Time

1


Complexity isn't the same as efficiency.

- A good $O(n^2)$ implementation can be better than a bad $O(n)$.
 - For a while.
- Eventually, as n increases, $O(n)$ will always outperform $O(n^2)$ etc.

```
def n_sum(sequence):  
    total = 0  
    for i in range(len(sequence)):  
        total += sequence[i]  
        time.sleep(0.001)  
    return total
```

lec_fast_slow_functions.py

```
def n2_sum(sequence):  
    total = 0  
    for i in range(len(sequence)):  
        counter = 0  
        while counter < i:  
            counter += 1  
        total += sequence[counter]  
  
    return total
```



Why do I care?

■ Everyone

- Without $O()$ notation we can't discuss how algorithms compare.
- Without $O()$ can't discuss why some tasks are effectively impossible (travelling salesman).
- You should be trying to write good, efficient code. Profiling helps you do this.
- Ethical Hackers - $O()$ important in discussing password security.
- Games Tech - $O()$ explains the need for path finding and graphics workarounds.

Recap

Profiling help determines the actual performance of your code.

- Statistical profilers.
 - Accurate-ish
- Instrumental profilers.
 - Insert additional instructions.
 - Accurate but slows things down.

$O()$ describes algorithm complexity.

- Time/space.
- How your code should scale.
 - Lots of real world issues can mess it up.
 - Memory limits etc.

● $O(1) < O(\log n) < O(n) < O(n \log n) < O(n^2) < O(2^n) < O(n!)$

● $\geq O(2^n)$ means exponential.

● $< O(n^2)$ means polynomial.

Profiling

Efficiency

Optimization

Profilers

$O()$ notation

Simple algorithms

Good algorithms

Bad algorithms

Recap

The End